

Neural Networks Basics

CSE 4237 - Soft Computing

Mir Tafseer Nayeem
Faculty Member, CSE AUST
tafseer.nayeem@gmail.com

Binary Classification




→ **1** (Cat) vs **0** (non-Cat)

Example: Cat vs Non-Cat


- The goal is to **train a classifier** with an input image represented by a **feature vector x** .
- To predict whether the corresponding label y is 1 or 0.
- In this case, whether this is a cat image (1) or a non-cat image (0).

Binary Classification




- An image is stored in the computer in three separate matrices corresponding to the **Red**, **Green**, and **Blue** color channels of the image.
- The three matrices have the **same size as the image**, for example, the resolution of the cat image is **64 pixels X 64 pixels**, the **three matrices (RGB) are 64 X 64 each**.

Binary Classification



- The value in a cell represents the pixel intensity which will be used to create a feature vector of n dimension. In pattern recognition and machine learning, a feature vector represents an object, in this case, a cat or no cat.

Binary Classification



- To create a feature vector, x , the pixel intensity values will be “unroll” or “reshape” for each color. The dimension of the input feature vector x is $n_x = 64 \times 64 \times 3 = 12\,288$.

Logistic Regression


- Logistic regression is a learning algorithm used in a supervised learning problem when the output y are all either zero or one.
- The goal of logistic regression is to minimize the error between its predictions and training data.
- Given an image represented by a feature vector x , the algorithm will evaluate the probability of a cat being in that image.

$$\text{Given } x, \hat{y} = P(y = 1|x), \text{ where } 0 \leq \hat{y} \leq 1$$

Logistic Regression

The parameters used in Logistic regression are:

- The input features vector: $x \in \mathbb{R}^{n_x}$, where n_x is the number of features
- The training label: $y \in 0,1$
- The weights: $w \in \mathbb{R}^{n_x}$, where n_x is the number of features
- The bias $b \in \mathbb{R}$
- The output: $\hat{y} = \sigma(w^T x + b)$
- Sigmoid function: $s = \sigma(w^T x + b) = \sigma(z) = \frac{1}{1 + e^{-z}}$




Parameters

Logistic Regression : Role of bias (b)

- The bias value **allows the activation function to be shifted to the left or right**, to better fit the data.
- Changes to the weights **alter the steepness of the sigmoid curve**, whilst the bias offsets it, shifting the entire curve so it fits better.
- **Bias only influences the output values**, it doesn't interact with the actual input data. That's why it is called bias.
- You can think of the bias as a measure of **how easy it is to get a node to fire**.
 - **For a node with a large bias**, the output will tend to be intrinsically high, with small positive weights and inputs producing large positive outputs (near to 1).
 - **Biases can be also negative**, leading to sigmoid outputs near to 0.
 - **If the bias is very small (or 0)**, the output will be decided by the values of weights and inputs alone.

Logistic Regression


- $(wTx + b)$ is a linear function like $(ax + b)$, but since we are looking for a probability constraint between $[0,1]$, the sigmoid function is used.
- The function is bounded between $[0,1]$ as shown in the graph below.



Logistic Regression

Some observations from the graph:

- If z is a large positive number, then $\sigma(z) = 1$
- If z is small or large negative number, then $\sigma(z) = 0$
- If $z = 0$, then $\sigma(z) = 0.5$



Logistic Regression: Cost Function

- To train the parameters w and b , we need to define a cost function.

$$\hat{y}^{(i)} = \sigma(w^T x^{(i)} + b), \text{ where } \sigma(z^{(i)}) = \frac{1}{1 + e^{-z^{(i)}}}$$

Given $\{(x^{(1)}, y^{(1)}), \dots, (x^{(m)}, y^{(m)})\}$, we want $\hat{y}^{(i)} \approx y^{(i)}$

$x^{(i)}$ the i -th training example

Loss (error) function:

- Loss function measures the discrepancy between the prediction ($\hat{y}^{(i)}$) and the desired output ($y^{(i)}$).
- In other words, the loss function computes the error for a single training example.


Error / Loss Function

Squared Error Function: $L(\hat{y}^{(i)}, y^{(i)}) = \frac{1}{2}(\hat{y}^{(i)} - y^{(i)})^2$

- We can see an extra (1/2) in the right side of the equation. Does it matter?
- It is because when you take the derivative of the cost function, that is used in updating the parameters during gradient descent, that 2 in the power get cancelled with the (1/2) multiplier.
- These techniques are or somewhat similar are widely used in math in order **"To make the derivations mathematically more convenient"**.

Is squared error function a good choice?

- The squared error function (commonly used function for linear regression) is not very suitable for logistic regression.
 - In case of logistic regression, the hypothesis / prediction is non-linear (sigmoid function), which makes the square error function to be non-convex.
 - On the other hand, logarithmic function is a convex function for which there is no local optima, so gradient descent works well.
- If you are doing binary classification, squared error function generally also penalize examples that **are correctly classified but are still near the decision boundary**, thus creating a "**margin.**"
- Gradient descent waste a lot of time getting predictions very close to $\{0, 1\}$



Logistic Regression: Cross Entropy Loss

$$L(\hat{y}^{(i)}, y^{(i)}) = -(y^{(i)} \log(\hat{y}^{(i)}) + (1 - y^{(i)}) \log(1 - \hat{y}^{(i)}))$$


- If $y^{(i)} = 1$: $L(\hat{y}^{(i)}, y^{(i)}) = -\log(\hat{y}^{(i)})$ where $\log(\hat{y}^{(i)})$ and $\hat{y}^{(i)}$ should be close to 1
- If $y^{(i)} = 0$: $L(\hat{y}^{(i)}, y^{(i)}) = -\log(1 - \hat{y}^{(i)})$ where $\log(1 - \hat{y}^{(i)})$ and $\hat{y}^{(i)}$ should be close to 0
- **Cost function**
 - The cost function is the average of the loss function of the entire training set. We are going to find the parameters w and b that minimize the overall cost function.

$$J(w, b) = \frac{1}{m} \sum_{i=1}^m L(\hat{y}^{(i)}, y^{(i)}) = -\frac{1}{m} \sum_{i=1}^m [(y^{(i)} \log(\hat{y}^{(i)}) + (1 - y^{(i)}) \log(1 - \hat{y}^{(i)}))]$$

Gradient Descent

$$\text{Recap: } \hat{y} = \sigma(w^T x + b), \quad \sigma(z) = \frac{1}{1+e^{-z}}$$

$$J(w, b) = \frac{1}{m} \sum_{i=1}^m \mathcal{L}(\hat{y}^{(i)}, y^{(i)}) = -\frac{1}{m} \sum_{i=1}^m y^{(i)} \log \hat{y}^{(i)} + (1 - y^{(i)}) \log(1 - \hat{y}^{(i)})$$




We want to find parameters **W, b** that minimize **J(W, b)**

Gradient Descent

- Our cost function is convex.
- First we initialize **w** and **b** to **0,0** or initialize them to a **random value** in the convex function and then try to improve the values the reach minimum value.
- In Logistic regression people **always use 0,0 instead of random**.
- **This function is convex**, no matter where you initialize you should get to the **global optimal point or roughly close the global optimal point**.


We want to parameters **W, b** that minimize **J(W, b)**




Gradient Descent

- Gradient starts at the initial point and take a step in the steepest downhill direction after each iteration.
- It will try to reach to the global optimum or somewhere near to the global optimum.

We want to parameters $\mathbf{W}$, $\mathbf{b}$ that minimize $\mathbf{J}(\mathbf{W}, \mathbf{b})$



Gradient Descent



Ignore $\mathbf{b}$ for now to make it a one dimensional plot rather than a higher dimensional plot.

- α = **Learning Rate**: How bigger step we choose at each iteration of gradient descent.
- Definition of a derivative:
 - Slope of a function at a point.

// Repeatedly do that until the algorithm converges.

Repeat {


Learning Rate

$$w := w - \alpha \frac{dJ(w)}{dw}$$

Update or change you want to make to the parameter w

}

$$\begin{aligned}\text{Slope} &= \frac{\text{Change in Y}}{\text{Change in X}} \\ &= \frac{\Delta y}{\Delta x}\end{aligned}$$



Gradient Descent : Actual Update Rule

We want to parameters
W, b that minimize **J(W, b)**


Partial Derivative

$J(w, b)$


$$w := w - \alpha \frac{\partial J(w, b)}{\partial w}$$

$$b := b - \alpha \frac{\partial J(w, b)}{\partial b}$$

Derivatives : Intuition



- $a = 2$ $f(a) = 6$
 $a = 2.001$ $f(a) = 6.003$

Slope (derivative) of $f(a)$ at $a = 2$ is 3

- $a = 5$ $f(a) = 15$
 $a = 5.001$ $f(a) = 15.003$

Slope (derivative) of $f(a)$ at $a = 5$ is also 3

$$\frac{df(a)}{da} = 3$$


What does $\frac{d}{dx}x^2 = 2x$ mean?

The slope or "rate of change" at any point is **2x**.

If we shift **a** by 0.001 then **f(a)** shift by 3 times 0.001.

Do we actually need Gradient Descent?


- Let's pretend that we only have 1 weight. To find the ideal value of our weight that will minimize our cost, we need to try a bunch of values for W , let's say we test 1000 values. That doesn't seem so bad, after all, my computer is pretty fast.



- It takes about **0.04 seconds** to check **1000 different weight** values for our neural network.
- Since we've computed the cost for a wide range values of W , we can **just pick the one with the smallest cost**.


Do we actually need Gradient Descent?

- Let's next consider 2 weights for a moment. To maintain the same precision we now need to check **1000 times 1000, or one million values**. This is a lot of work, even for a fast computer.



- After our 1 million evaluations we've found our solution, but it took an **agonizing 40 seconds!** Searching through **three weights would take a billion evaluations, or 11 hours!**
- Searching through all 9 weights we need for our simple network would take **1,268,391,679,350,583.5 years**. (Over a quadrillion years). So for that reason, the **"just try everything"** or **brute force optimization method** is clearly not going to work.

A Famous Quote



Computation Graph


- Neural Networks are organized in terms of a forward pass or backward pass.
- Forward Pass / Propagation
 - Which we compute the output of the neural network
- Backward Pass / Propagation
 - Which we use to compute gradients / derivatives
- **Computation Graph**
 - Explains why it is organized in this way.

Computation Graph

$$J(a, b, c) = 3(a + bc)$$

3 steps of computation:

1. $u = bc$
2. $v = a + u$
3. $j = 3v$




Logistic regression : Forward Propagation


$$z = w^T x + b$$

$$\hat{y} = a = \sigma(z)$$

$$\mathcal{L}(a, y) = -(y \log(a) + (1 - y) \log(1 - a))$$



Logistic regression : Backward Propagation



Rules for derivatives of logarithmic expressions

$$\mathcal{L}(a, y) = -(y \log(a) + (1 - y) \log(1 - a))$$

$$\frac{\delta}{\delta x} \log(\text{expression}) = \frac{1}{\text{expression}} \cdot \frac{\delta}{\delta x} \text{expression}$$

Examples:

$$\frac{\delta}{\delta x} \log(x) = \frac{1}{x} \cdot \frac{\delta}{\delta x} x = \frac{1}{x} \cdot 1 = \frac{1}{x}$$

If you are unsure about your derivative check this [link](#) to generate the derivation steps.

Logistic regression : Backward Propagation

$$\mathcal{L}(a, y) = -(y \log(a) + (1 - y) \log(1 - a))$$

Ignoring the (-) sign for now.

$$\begin{aligned} & \frac{d}{da} [y \ln(a) + (1 - y) \ln(1 - a)] \\ &= y \cdot \frac{d}{da} [\ln(a)] + (1 - y) \cdot \frac{d}{da} [\ln(1 - a)] \\ &= y \cdot \frac{1}{a} + (1 - y) \cdot \frac{1}{1 - a} \cdot \frac{d}{da} [1 - a] \\ &= \frac{y}{a} + \frac{(1 - y) \left(\frac{d}{da} [1] - \frac{d}{da} [a] \right)}{1 - a} \end{aligned}$$

log (x) refers to **e base log or the natural logarithm (ln(x))** in mathematical analysis, physics, chemistry, statistics, economics, and some engineering fields.

Logistic regression : Backward Propagation

$$= \frac{y}{a} + \frac{(1 - y) \left(\frac{d}{da}[1] - \frac{d}{da}[a] \right)}{1 - a}$$

$$= \frac{y}{a} + \frac{(1 - y)(0 - 1)}{1 - a}$$

$$= \frac{y}{a} + \frac{y - 1}{1 - a}$$

$$= \frac{y}{a} - \frac{1 - y}{1 - a}$$

Finally, adding the (-) sign.


$$= -\frac{y}{a} + \frac{1 - y}{1 - a}$$

Logistic regression : Backward Propagation

$$z = w^T x + b$$

$$\hat{y} = a = \sigma(z)$$

$$\mathcal{L}(a, y) = -(y \log(a) + (1 - y) \log(1 - a))$$



Applying Chain Rule

$$= \frac{\partial \mathcal{L}}{\partial a} \cdot \frac{\partial a}{\partial z}$$

$$= -\frac{y}{a} + \frac{1 - y}{1 - a}$$


Logistic regression : Backward Propagation

$$\begin{aligned}\frac{\partial a}{\partial z} &= \frac{\partial}{\partial z} \sigma(z) \\ &= \frac{\partial}{\partial z} \left[\frac{1}{1 + e^{-z}} \right] \\ &= \frac{\partial}{\partial z} (1 + e^{-z})^{-1} \\ &= -(1 + e^{-z})^{-2} (-e^{-z}) \\ &= \frac{e^{-z}}{(1 + e^{-z})^2} \\ &= \frac{1}{1 + e^{-z}} \cdot \frac{e^{-z}}{1 + e^{-z}}\end{aligned}$$

$$a = \sigma(z) = \frac{1}{1 + e^{-z}}$$

$$\begin{aligned}&= \frac{1}{1 + e^{-z}} \cdot \frac{(1 + e^{-z}) - 1}{1 + e^{-z}} \\ &= \frac{1}{1 + e^{-z}} \cdot \left(\frac{1 + e^{-z}}{1 + e^{-z}} - \frac{1}{1 + e^{-z}} \right) \\ &= \frac{1}{1 + e^{-z}} \cdot \left(1 - \frac{1}{1 + e^{-z}} \right) \\ &= \sigma(z) \cdot (1 - \sigma(z)) \\ &= a \cdot (1 - a)\end{aligned}$$

Logistic regression : Backward Propagation



Applying Chain Rule

$$= \frac{\partial \mathcal{L}}{\partial a} \cdot \frac{\partial a}{\partial z}$$

$$= -\frac{y}{a} + \frac{1-y}{1-a}$$


$$= \left(-\frac{y}{a} + \frac{1-y}{1-a}\right) * (a * (1-a))$$

$$= -y * (1-a) + a * (1-y)$$

$$= -y + ay + a - ay$$

$$= a - y$$

Logistic regression : Backward Propagation



$$dw_1 = \frac{\partial \mathcal{L}(a, y)}{\partial w_1} = \frac{\partial \mathcal{L}}{\partial w_1}$$

$$\frac{\partial \mathcal{L}}{\partial z} = \frac{\partial \mathcal{L}}{\partial a} \cdot \frac{\partial a}{\partial z}$$

$$dz = a - y$$


Applying
Chain Rule

$$= \frac{\partial \mathcal{L}}{\partial z} \cdot \frac{\partial z}{\partial w_1}$$

$$= dz * \frac{\partial}{\partial w_1} (w_1x_1 + w_2x_2 + b)$$

$$dw_1 = dz * x_1$$

Logistic regression : Backward Propagation



$$dw_2 = \frac{\partial \mathcal{L}(a, y)}{\partial w_2} = \frac{\partial \mathcal{L}}{\partial w_2}$$

$$\frac{\partial \mathcal{L}}{\partial z} = \frac{\partial \mathcal{L}}{\partial a} \cdot \frac{\partial a}{\partial z}$$

$$dz = a - y$$


Applying
Chain Rule

$$= \frac{\partial \mathcal{L}}{\partial z} \cdot \frac{\partial z}{\partial w_2}$$

$$= dz * \frac{\partial}{\partial w_2} (w_1x_1 + w_2x_2 + b)$$

$$dw_2 = dz * x_2$$

Logistic regression : Backward Propagation



$$db = \frac{\partial \mathcal{L}(a, y)}{\partial b} = \frac{\partial \mathcal{L}}{\partial b}$$

$$\frac{\partial \mathcal{L}}{\partial z} = \frac{\partial \mathcal{L}}{\partial a} \cdot \frac{\partial a}{\partial z}$$

Applying
Chain Rule

$$= \frac{\partial \mathcal{L}}{\partial z} \cdot \frac{\partial z}{\partial b}$$

$$dz = a - y$$

$$= dz * \frac{\partial}{\partial b} (w_1x_1 + w_2x_2 + b)$$

$$db = dz$$

Updating the Parameters: w_1 , w_2 and b

$$w_1 := w_1 - \alpha * dw_1$$

$$w_2 := w_2 - \alpha * dw_2$$

$$b := b - \alpha * db$$

**This is one step of
Gradient Descent on a
single example.**

Learning Rate



Logistic regression Gradient descent on m examples

Basic Parameters

x_1	Feature
x_2	Feature
w_1	Weight of the first feature.
w_2	Weight of the second feature.
b	Logistic Regression parameter (Bias).
m	Number of training examples
$y(i)$	Expected output of i

Logistic regression Gradient descent on m examples

$$z^{(i)} = w^T x^{(i)} + b$$

$$\hat{y}^{(i)} = a^{(i)} = \text{sigmoid}(z^{(i)})$$

$$\mathcal{L}(a^{(i)}, y^{(i)}) = -y^{(i)} \log(a^{(i)}) - (1 - y^{(i)}) \log(1 - a^{(i)})$$

$$J = \frac{1}{m} \sum_{i=1}^m \mathcal{L}(a^{(i)}, y^{(i)})$$

$$\frac{\partial}{\partial w_1} J(w, b) = \frac{1}{m} \sum_{i=1}^m \underbrace{\frac{\partial}{\partial w_1} \mathcal{L}(a^{(i)}, y^{(i)})}_{dw_1^{(i)}} \quad \text{For the example } (x^{(i)}, y^{(i)})$$

Logistic regression Gradient descent on m examples

Derivatives: All it turned out as simple arithmetic operations

$d(a)$	$-(y/a) + ((1-y) / (1-a))$
$d(z)$	$a - y$
$d(w1)$	$x1 * d(z)$
$d(w2)$	$x2 * d(z)$
$d(b)$	$d(z)$

Logistic regression Gradient descent on m examples

$J = 0$; $dw1 = 0$; $dw2 = 0$; $db = 0$;

$w1 = 0$; $w2 = 0$; $b = 0$;

for $i = 1$ to m

Forward pass

$z(i) = w1 * x1(i) + w2 * x2(i) + b$

$a(i) = \text{sigmoid}(z(i))$

$J += (y(i) * \log(a(i)) + (1 - y(i)) * \log(1 - a(i)))$

Backward pass

$dz(i) = a(i) - y(i)$

$dw1 += dz(i) * x1(i)$

$dw2 += dz(i) * x2(i)$

$db += dz(i)$

$\updownarrow$
 $n = 2$

$J /= m$

$dw1 /= m$

$dw2 /= m$

$db /= m$

Gradient descent

$w1 = w1 - \alpha * dw1$

$w2 = w2 - \alpha * dw2$

$b = b - \alpha * db$

w1, w2, b are the accumulators and single instances for the all **m** training examples.

One iteration of gradient descent

Logistic regression Gradient descent on m examples

- Previous slide is **just one step of Gradient Descent**, we need to repeat it **multiple times** in order to take multiple steps of gradient descent.
- There are **weaknesses** in the previous implementation. In order to implement we need to **write two for loops**.
- Having explicit for loops in your code make your **code less efficient**.
- Solution:- **vectorization techniques**
- To train with larger datasets we need to take the help from vectorization techniques **without using for loops**.

LR Gradient descent on m examples (modified)

```
J = 0; dw1 = 0; dw2 = 0; db = 0;  
w1 = 0; w2 = 0; b = 0;
```

```
dw = np.zeros ((nx, 1))
```

for i = 1 to m

Forward pass

```
z(i) = w1*x1(i) + w2*x2(i) + b
```

```
a(i) = sigmoid(z(i))
```

```
J += - (y(i)*log(a(i)) + (1-y(i))*log(1-a(i)))
```

Backward pass

```
dz(i) = a(i) - y(i)
```

```
dw1 += dz(i) * x1(i)
```

```
dw2 += dz(i) * x2(i)
```

```
db += dz(i)
```

n = 2

```
dw += x(i) * dz(i)
```

Logistic regression Gradient descent on m examples

$$J = J / m$$

$$\begin{aligned} dw1 &= dw1 / m \\ dw2 &= dw2 / m \end{aligned}$$

$$db = db / m$$

$$dw = dw / m$$

Gradient descent

$$w1 = w1 - \alpha * dw1$$

$$w2 = w2 - \alpha * dw2$$

$$b = b - \alpha * db$$

w1, w2, b are the accumulators and single instances for the all **m** training examples.

We have gone from 2 for loops to 1 for loop, we still have one for loop that loops over individual training examples.

Vectorizing Logistic Regression (Forward)

$$z^{(1)} = w^T x^{(1)} + b$$

$$a^{(1)} = \sigma(z^{(1)})$$

1st training example

$$z^{(2)} = w^T x^{(2)} + b$$

$$a^{(2)} = \sigma(z^{(2)})$$

2nd training example

$$z^{(3)} = w^T x^{(3)} + b$$

$$a^{(3)} = \sigma(z^{(3)})$$

3rd training example

We need to do it m times if you have **m training examples**.

$$\mathbf{X} = \begin{bmatrix} \vdots & \vdots & \vdots \\ x^{(1)} & x^{(2)} & \dots & x^{(m)} \\ \vdots & \vdots & \vdots \end{bmatrix} \text{ while } \mathbb{R}^{n_x \times m}$$

$$\begin{bmatrix} w^T \end{bmatrix} \begin{bmatrix} \vdots & \vdots & \vdots \\ x^{(1)} & x^{(2)} & \dots & x^{(m)} \\ \vdots & \vdots & \vdots \end{bmatrix}$$

$$[z^{(1)} \ z^{(2)} \ z^{(3)} \ \dots \ z^{(m)}] = w^T X + [b \ b \ \dots \ b] = [w^T x^{(1)} + b \quad w^T x^{(2)} + b \quad \dots \quad w^T x^{(m)} + b]$$

→ (1 x m) dimension
→ (1 x m)

Vectorizing Logistic Regression (Forward)

$$[z^{(1)} \ z^{(2)} \ z^{(3)} \ \dots \ z^{(m)}] = w^T X + [b \ b \ \dots \ b] = [w^T x^{(1)} + b \quad w^T x^{(2)} + b \quad \dots \quad w^T x^{(m)} + b]$$

$$Z = [z^{(1)} \ z^{(2)} \ z^{(3)} \ \dots \ z^{(m)}]$$

(1 x m) dimension

Broadcasting

$$Z = np.dot(w.T, X) + b$$

(1, 1) dimension

$$A = [a^{(1)} \ a^{(2)} \ a^{(3)} \ \dots \ a^{(m)}] = \sigma(Z)$$

Gradient Computation

$$dz^{(1)} = a^{(1)} - y^{(1)} \quad dz^{(2)} = a^{(2)} - y^{(2)} \quad dz^{(3)} = a^{(3)} - y^{(3)}$$
$$dZ = [dz^{(1)} dz^{(2)} \dots dz^{(m)}]$$

$$A = [a^{(1)} a^{(2)} a^{(3)} \dots a^{(m)}] \quad Y = [y^{(1)} y^{(2)} y^{(3)} \dots y^{(m)}]$$
$$dZ = A - Y = [(a^{(1)} - y^{(1)})(a^{(2)} - y^{(2)})(a^{(3)} - y^{(3)}) \dots (a^{(m)} - y^{(m)})]$$

Gradient Computation

$$\begin{aligned}
 dw &= 0 \\
 dw &+ = X^{(1)} dz^{(1)} \\
 dw &+ = X^{(2)} dz^{(2)} \\
 &\vdots \\
 dw &+ = X^{(m)} dz^{(m)} \\
 dw &/ = m
 \end{aligned}$$

$$\begin{aligned}
 db &= \frac{1}{m} \sum_{i=1}^m dz^{(i)} \\
 db &= \frac{1}{m} np.sum(dZ)
 \end{aligned}$$

$$dw = \frac{1}{m} X \cdot dZ^T$$

$$\begin{aligned}
 db &= 0 \\
 db &+ = dz^{(1)} \\
 db &+ = dz^{(2)} \\
 &\vdots \\
 db &+ = dz^{(m)} \\
 db &/ = m
 \end{aligned}$$

$$dw = \frac{1}{m} \begin{bmatrix} \vdots & \vdots & & \vdots \\ x^{(1)} & x^{(2)} & \dots & x^{(m)} \\ \vdots & \vdots & & \vdots \end{bmatrix} \begin{pmatrix} dz^{(1)} \\ dz^{(2)} \\ \vdots \\ dz^{(m)} \end{pmatrix}$$

$$dw = \frac{1}{m} [X^{(1)} dz^{(1)} + X^{(2)} dz^{(2)} + \dots + X^{(m)} dz^{(m)}]$$

Implementing Logistic Regression

$J = 0, dw_1 = 0, dw_2 = 0, db = 0$

for $i = 1$ to m :

$$z^{(i)} = w^T x^{(i)} + b$$

$$a^{(i)} = \sigma(z^{(i)})$$

$$J += -[y^{(i)} \log a^{(i)} + (1 - y^{(i)}) \log(1 - a^{(i)})]$$

$$dz^{(i)} = a^{(i)} - y^{(i)}$$

$$dw_1 += x_1^{(i)} dz^{(i)}$$

$$dw_2 += x_2^{(i)} dz^{(i)}$$

$$db += dz^{(i)}$$

$$J = J/m, dw_1 = dw_1/m, dw_2 = dw_2/m$$

$$db = db/m$$

Single Iteration of Gradient Descent

$$Z = w^T X + b$$

$$Z = np.dot(w, X) + b$$

$$A = \sigma(Z)$$

$$dZ = A - Y$$

$$dw = \frac{1}{m} X \cdot dZ^T$$

$$db = \frac{1}{m} \sum_{i=1}^m dz^{(i)}$$


$$db = \frac{1}{m} np.sum(dZ)$$

$$w := w - \alpha * dw$$

$$b := b - \alpha * db$$

Gradient
Update

What does this have to do with the brain?



END